When Short Responses Are Hard to Distinguish in Translation Audits

João Victor M. Correa[1,2] **Pedro Shiozawa**[3] **Victor Hierro Moraes Ferreira**[4] **Emanoel Cristhian Leão Marinho**[4]

[1]Universidade de São Paulo [2]Great People Mental Health
[3]Faculdade de Ciências Médicas da Santa Casa de São Paulo [4]Universidade Federal do Pará

MultiPsyche 2026 | AACL-IJCNLP 2026

Introduction

Workplace surveys often contain brief or repeated answers. Translation makes these responses accessible to English-centered analysis tools, including resources used in mental-health research.

We examine what reference-free audit tools reveal about Portuguese-to-English survey translations and how response length, repetition and lexical coverage shape their interpretation.

Data and methods

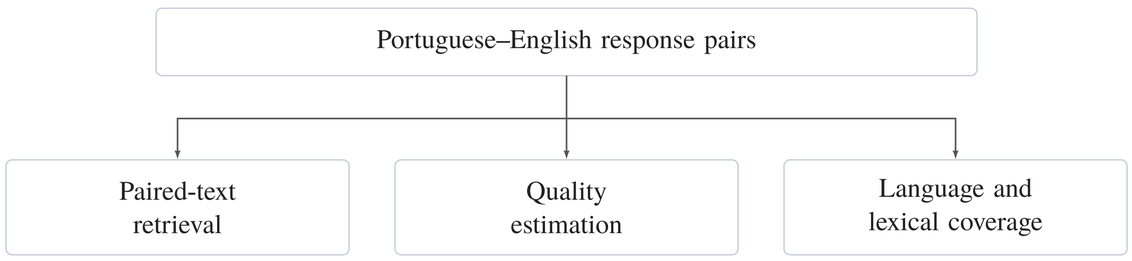
The corpus contains **44,118 response pairs** from **14,706 survey-specific records** across four Brazilian workplace survey cohorts from 2023 to 2025. Each record contributes three answers. One third of pairs come from mental-health or well-being prompts.

Paired-text retrieval. Qwen3-Embedding-4B [1] ranks all 44,118 candidates in the other language. The paired target or an exact normalized duplicate counts as a correct match. We also tested the 8B encoder.

Quality estimation. COMETKiwi DA-XL [2] estimates translation quality from the Portuguese source and English translation.

Language and affective coverage. Lingua provides language flags. EmoAtlas [3] measures emotion-lexicon coverage in each language.

Retrieval, quality estimation and coverage are interpreted separately.



Interpretation

Retrieval measures how distinguishable a paired text is within this corpus. A missed match can reflect brief or repeated wording.

COMETKiwi follows a different pattern across response lengths. Reporting the measures separately keeps their meanings clear. Bilingual human ratings would calibrate the quality estimates.

Limitations

Results describe one language pair and one workplace-survey setting. Translation provenance is incomplete, and encoder comparisons remain within the Qwen3 family.

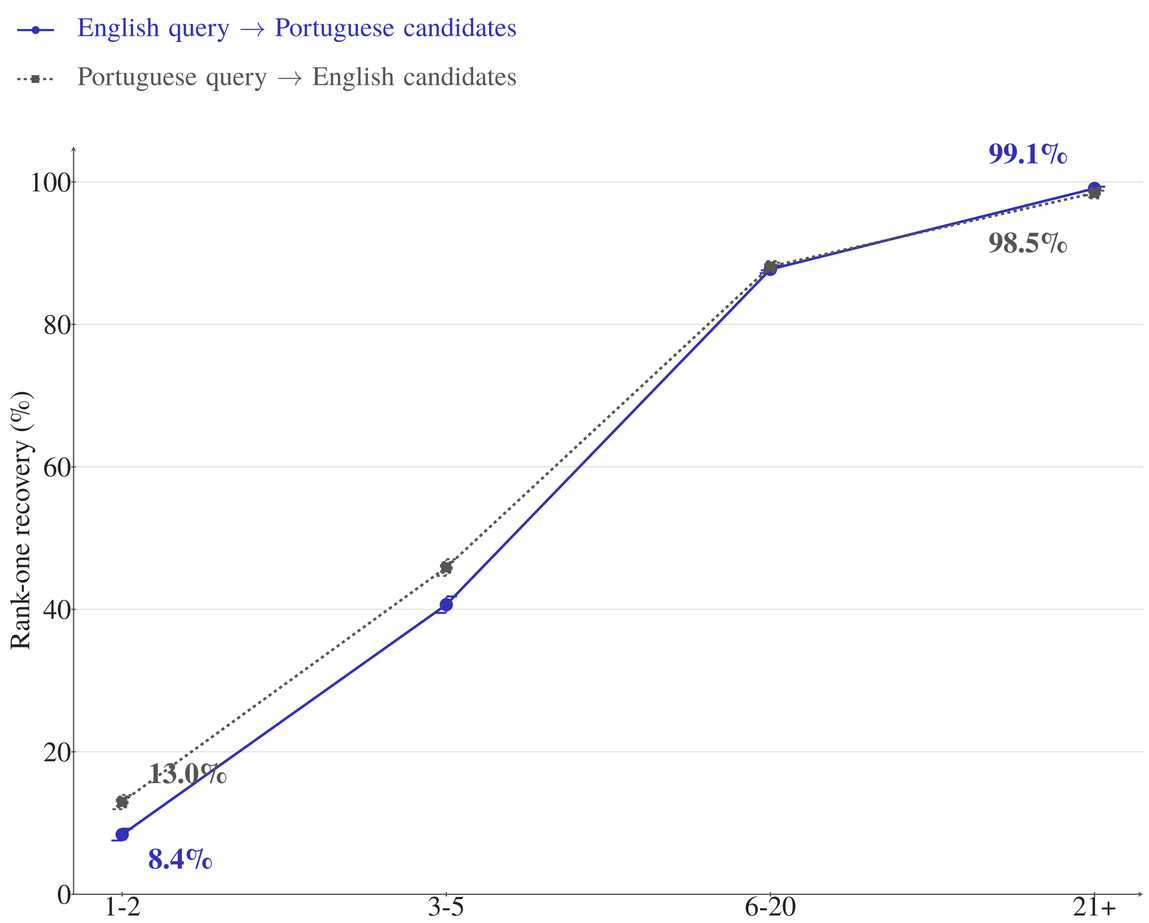
Assessing translation accuracy and emotional or cultural equivalence requires human assessment and validation in the target domain.

Survey comments remain confidential. The public code includes frozen configurations and synthetic tests.

Results

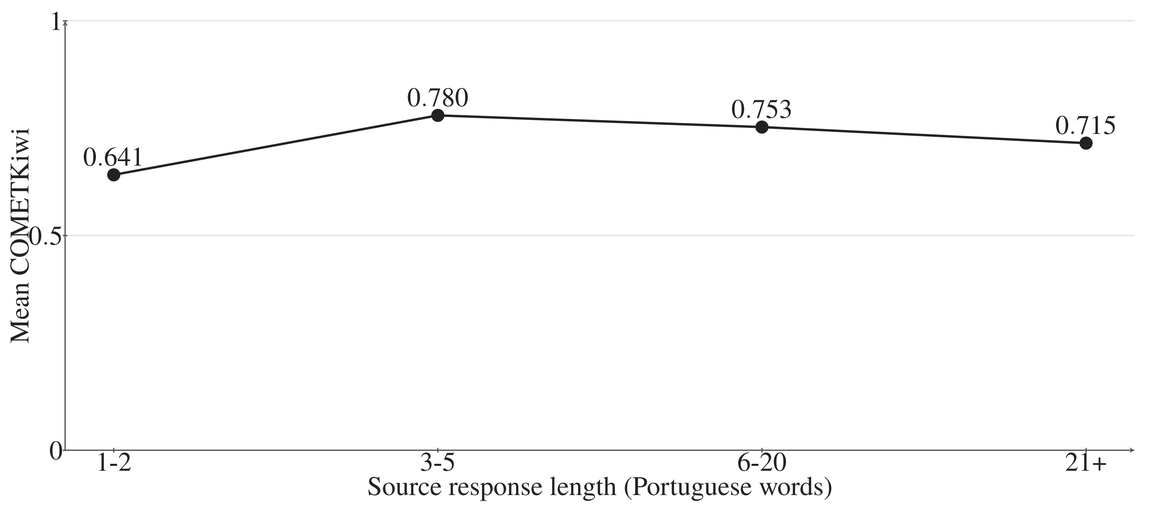
Retrieval by response length

Rank-one recovery finds the paired target or an exact normalized duplicate in the first position.



Error bars show 95% bootstrap intervals from 2,000 resamples, clustered by survey-specific record. Every query searches the full set of 44,118 candidates.

Quality estimates differ from retrieval



COMETKiwi scores are model estimates. Calibration against bilingual ratings is needed for this domain.

Responses with zero emotion-lexicon coverage

Portuguese	24.67%
English	23.48%

Zero coverage limits lexical comparisons across languages.

Selected references

[1] Zhang et al. (2025). Qwen3 Embedding. arXiv:2506.05176.
[2] Rei et al. (2023). Scaling up CometKiwi. WMT 2023.
[3] Semeraro et al. (2025). EmoAtlas. Behavior Research Methods.

Correspondence victormcorrea@usp.br

Code