When Short Responses Are Hard to Distinguish in Translation Audits

João Victor M. Correa[1,2] Pedro Shiozawa[3]
Victor Hierro Moraes Ferreira[4] Emanoel Cristhian Leão Marinho[4]
[1]Universidade de São Paulo [2]Great People Mental Health
[3]Faculdade de Ciências Médicas da Santa Casa de São Paulo
[4]Universidade Federal do Pará

MultiPsyche 2026 AACL-IJCNLP 2026

Translation audits for short survey responses



Response length **Repeated wording** **Lexicon coverage**

These conditions shape how the audit outputs can be interpreted.

An audit has to account for the information each answer contains.

Camera-ready article, Introduction and Data and Methods. Diagram summarizes the study setting.

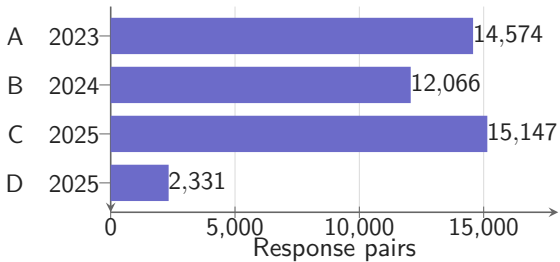
Four cohorts, three prompt families

44,118

Portuguese–English response pairs

14,706

survey-specific records



Positive workplace qualities

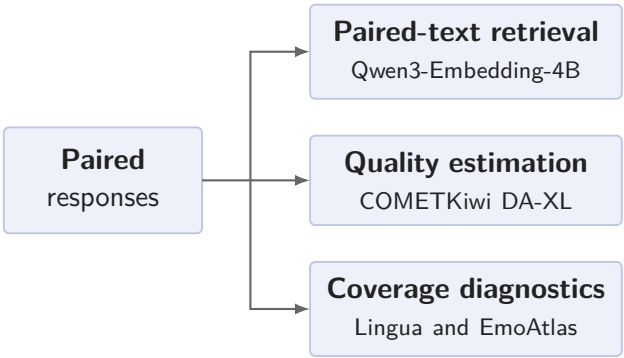
Areas for improvement

Mental health and well-being

Each retained record contributes one answer to each prompt family.

Camera-ready article, Table 1 and cohort-selection description. Counts represent survey-specific records.

Three audit views



Can the paired text be recovered?

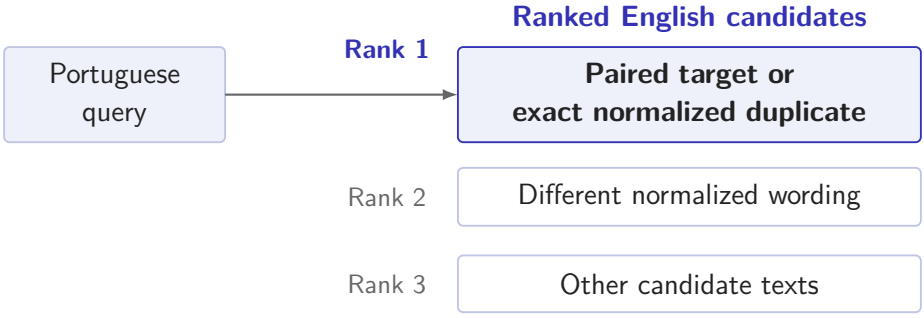
What quality does the model estimate?

Which language and emotion terms are recognized?

Retrieval, quality estimates and coverage are interpreted separately.

Camera-ready article, Figure 2 and Data and Methods. Model references appear on the final backup slide.

What rank-one recovery measures



71.80%

English query to Portuguese candidates

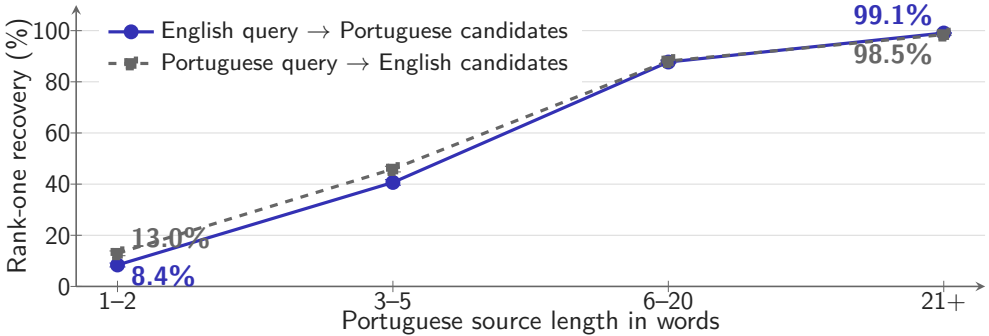
73.37%

Portuguese query to English candidates

Every query is ranked against all 44,118 texts in the other language.

Camera-ready article, retrieval definition and Table 2. Percentages are overall duplicate-aware R@1.

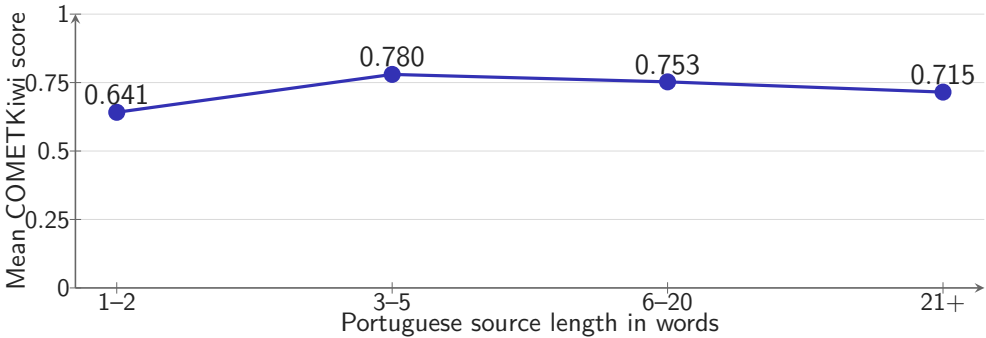
Retrieval rises sharply with response length



Short responses offer fewer cues for recovering their paired text.

Camera-ready article, Figure 1. 95% cluster-bootstrap intervals, 2,000 resamples over survey-specific records.

Quality estimates follow a different pattern

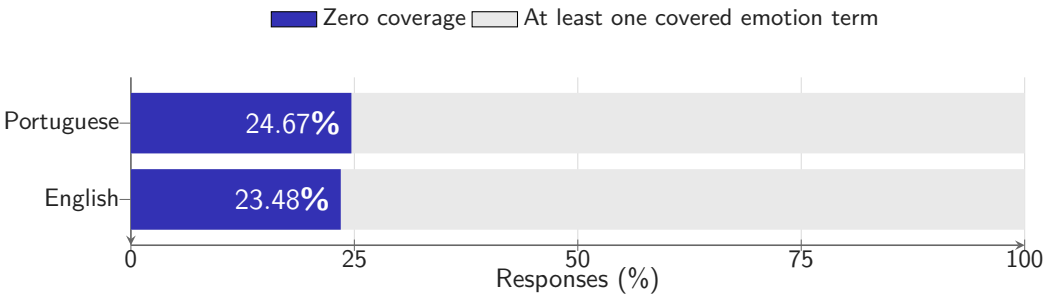


Bilingual human ratings would calibrate these quality estimates.

Camera-ready article, length-stratified COMETKiwi means. COMETKiwi DA-XL, Rei et al., WMT 2023.

Emotion-lexicon coverage

About one quarter of responses have **zero covered emotion terms**.



Affective comparisons need explicit reporting of lexical coverage.

Camera-ready article, Language and Affective Coverage. EmoAtlas, Semeraro et al., Behavior Research Methods, 2025.

Scope and validation

Study scope

Portuguese–English
Workplace surveys
Qwen3 4B and 8B encoders

Translation provenance is incomplete.
Response data remain confidential.

Next validation step

Bilingual human review
across lengths and prompts

↓

Translation quality
Emotional and
cultural meaning

Claims about translation quality and meaning require human assessment.

Camera-ready article, Discussion, Limitations and Ethics Statement. Survey comments remain confidential.

Report distinguishability, estimated quality and lexical coverage separately.

Length, repetition and resource coverage shape what an audit can tell us.

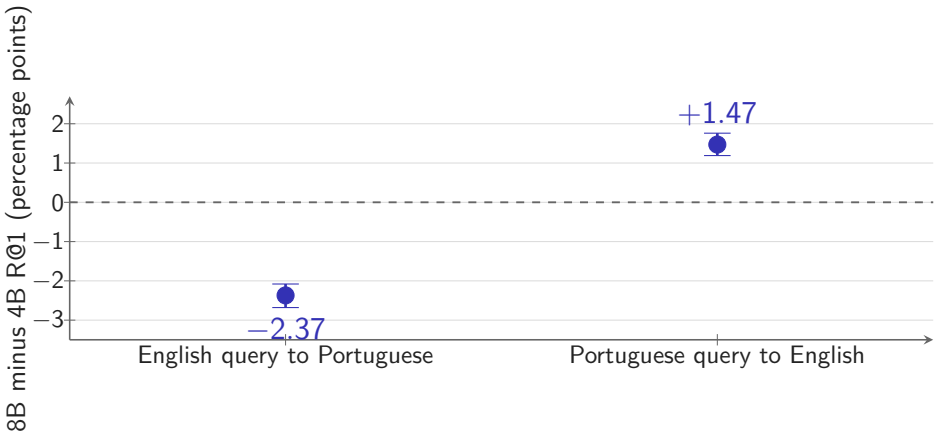
Code



João Victor M. Correa

victormcorrea@usp.br

Encoder sensitivity



The 8B encoder changes recovery in opposite directions.

Backup slide. Camera-ready article, Appendix C. Paired 8B minus 4B differences with 95% cluster-bootstrap intervals.

Selected references

Zhang et al., 2025
Qwen3 Embedding: Advancing Text Embedding and Reranking Through Foundation Models.
arXiv:2506.05176

Rei et al., 2023
Scaling Up CometKiwi: Unbabel-IST 2023 Submission for the Quality Estimation Shared Task.
Proceedings of WMT 2023

Semeraro et al., 2025
EmoAtlas: An Emotional Network Analyzer of Texts That Merges Psychological Lexicons, Artificial Intelligence, and Network Science.
Behavior Research Methods, DOI 10.3758/s13428-024-02553-7

Backup slide. Full references, Lingua version and model revisions appear in the camera-ready article and public code.